

An in-depth look at Slices and Maps

Session 07

Golang course by Exadel

24 Oct 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

Agenda

- ▶ Questions from the past session
- ▶ State of collections in other languages
- ▶ Maps in Go
- ▶ Slices in Go
- ▶ Next time...

Questions from the past session

Explain signed/unsigned integer type differences

✓ List of all available primitive types:

```
bool          // values are: 'true' or 'false'
string

// Numeric types:

uint          // either 32 or 64 bits depends on host platform
int           // same size as uint
uintptr       // an unsigned integer large enough to store the uninterpreted bits of a pointer value
uint8         // the set of all unsigned 8-bit integers (0 to 255)
uint16        // the set of all unsigned 16-bit integers (0 to 65535)
uint32        // the set of all unsigned 32-bit integers (0 to 4294967295)
uint64        // the set of all unsigned 64-bit integers (0 to 18446744073709551615)

int8          // the set of all signed 8-bit integers (-128 to 127)
int16         // the set of all signed 16-bit integers (-32768 to 32767)
int32         // the set of all signed 32-bit integers (-2147483648 to 2147483647)
int64         // the set of all signed 64-bit integers (-9223372036854775808 to 9223372036854775807)

float32       // the set of all IEEE-754 32-bit floating-point numbers
float64       // the set of all IEEE-754 64-bit floating-point numbers

complex64     // the set of all complex numbers with float32 real and imaginary parts
complex128    // the set of all complex numbers with float64 real and imaginary parts

byte          // alias for uint8
rune          // alias for int32 (represents a Unicode code point)
```

Binary representation (not only in Go)

```
1 package main
2
3 import "fmt"
4
5 func main() {
6     > var signed int8 = -126
7     > fmt.Printf(format: "Decimal Value: %+d ; Binary form: %b\n", signed, signed)
8     > // output:
9     > // Decimal Value: -126 ; Binary form: -1111110
10
11     > // the Most Significant Bit (MSB) == 1 in value 0b11111110
12     > // if the number is positive, this has a value of 0, but if the number is negative, the bit is 1.
13     > var unsigned uint8 = 0b11111110
14     > fmt.Printf(format: "Decimal Value: %+d ; Binary form: %b\n", unsigned, unsigned)
15     > // output:
16     > // Decimal Value: +254 ; Binary form: 11111110
17 }
```

▶ Check example at: `code/binary_representation_test.go`

▶ To check: "proposal: Go 2: use unsigned integers for all lengths" (<https://github.com/golang/go/issues/27460>)

More info about integers



▶ **Signed integers** are integers that can hold **positive and negative** numbers.

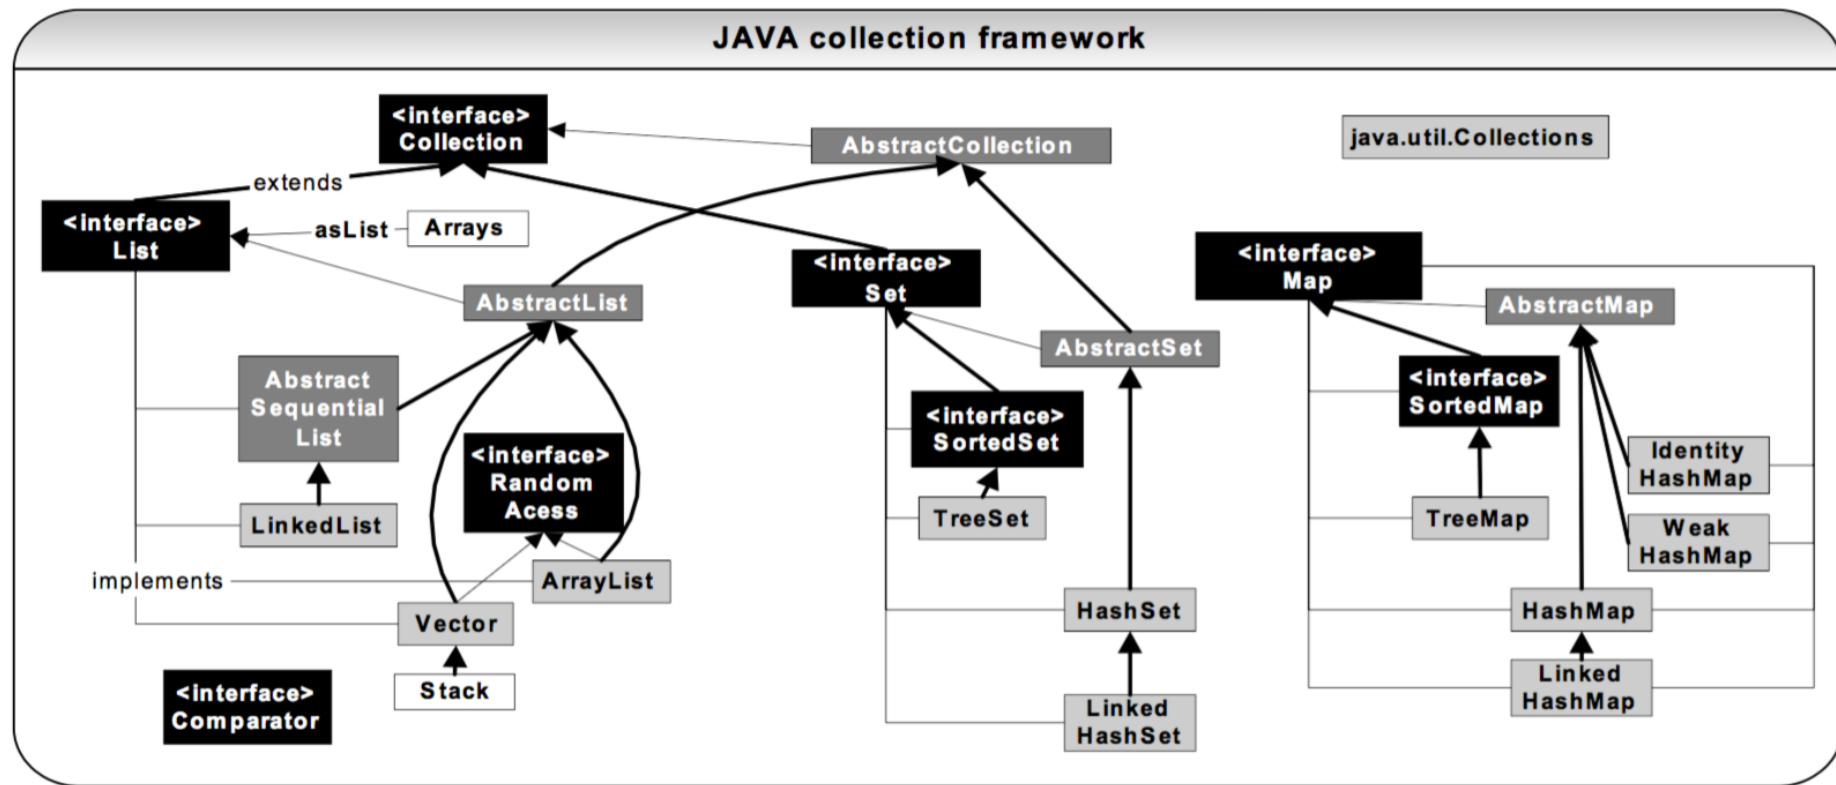
- for `int` on 32-bit CPU architecture it would be: from -2147483648 to 2147483647

▶ **Unsigned integers** are integers that can only hold **non-negative** whole numbers.

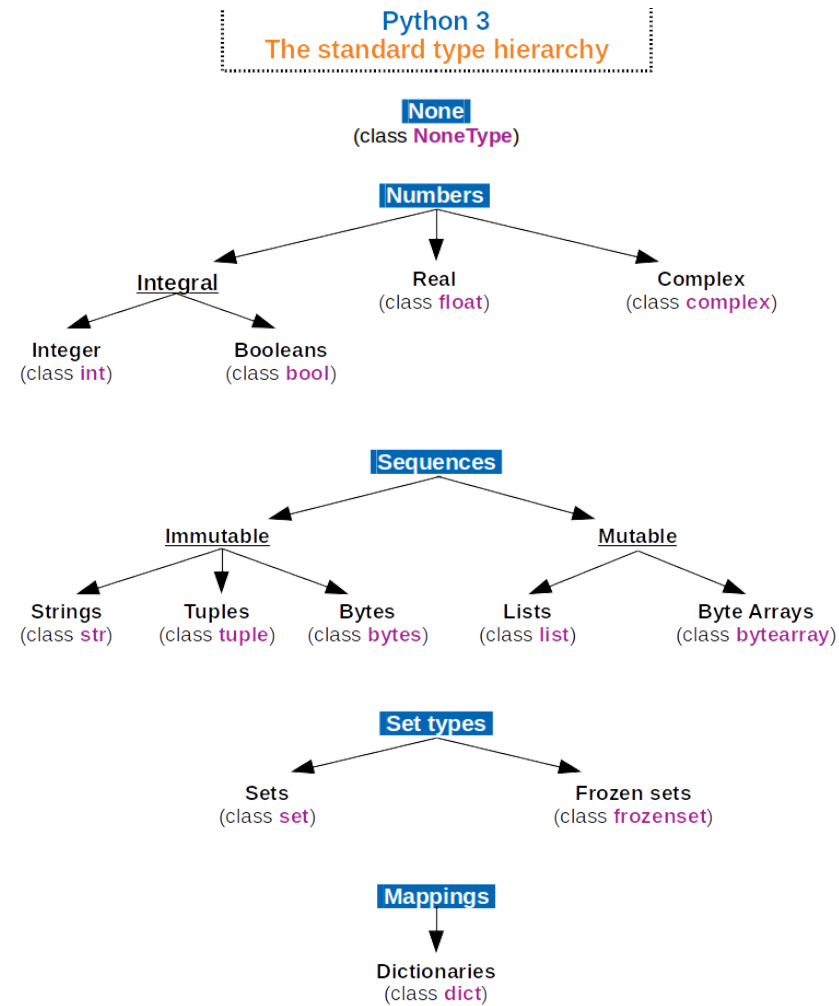
- for `uint` on 32-bit CPU architecture it would be: from 0 to 4294967295
- Consider the subtraction of two unsigned numbers, such as 3 and 5. 3 minus 5 is -2, but -2 can't be represented as an unsigned number.

State of collections in other languages

In Java (as of version 1.4)



In Python



Source: "Data type" by Wikipedia (https://en.wikipedia.org/wiki/Data_type)

Maps in Go

General information about Maps (1/2)

- ✓ Map is a language-supplied hash table.

`map[KeyType]ValueType`

- ✓ To create new map type we can use:
 - map literals;
 - `make()` function (as with slices).
 - The *initial capacity* does not bound its size:
 - maps grow to accommodate the number of items stored in them, with the exception of `nil` maps.
 - A `nil` map is equivalent to an empty map except that no elements may be added.
 - Might be used as “read-only” empty map (*safe to iterate over and retrieve values*).
 - Example:

```
make(map[string]int)
make(map[string]int, 100)
```

General information about Maps (2/2)

- ✓ **Keys** of map should support '*equality* ==' operator.
 - If <something> does not support '*equality*' it can not be used as a key.
 - Slices, maps, and functions can not be used as a keys!
 - In short, comparable types are:
 - boolean;
 - numeric;
 - string;
 - pointer;
 - channel;
 - interface types;
 - structs or arrays that contain only those types.
- ✓ Map's value type may be any type at all, including another map.
- ✓ Uninitialised value `var m map[string]int` is nil.
- ✓ Assigning a map copies the map reference, not the ~~map contents~~.

Maps under the hood

▶ As per Ian Taylor:

- "In the very early days what we call maps now were written as pointers, so you wrote `*map[int]int`. We moved away from that when we realized that no one ever wrote map without writing `*map`." (c) Ian Taylor
- More at: ["email thread"](https://groups.google.com/g/golang-nuts/c/SjuhSYDITm4/m/jnrp7rRxDQAJ) (https://groups.google.com/g/golang-nuts/c/SjuhSYDITm4/m/jnrp7rRxDQAJ)

```
115 // A header for a Go map.
116 type hmap struct {
117     > // Note: the format of the hmap is also encoded in cmd/compile/internal/reflectdata/reflect.go.
118     > // Make sure this stays in sync with the compiler's definition.
119     > count ..... int // # live cells == size of map. Must be first (used by len() builtin)
120     > flags ..... uint8
121     > B ..... uint8 // log_2 of # of buckets (can hold up to loadFactor * 2^B items)
122     > noverflow uint16 // approximate number of overflow buckets; see incrnoverflow for details
123     > hash0 .... uint32 // hash seed
124
125     > buckets ... unsafe.Pointer // array of 2^B Buckets. may be nil if count==0.
126     > oldbuckets unsafe.Pointer // previous bucket array of half the size, non-nil only when growing
127     > nevacuate  uintptr        // progress counter for evacuation (buckets less than this have been evacuated)
128
129     > extra *mapextra // optional fields
130 }
```

Map iteration order

✓ Map iteration order is not specified:

```
for key, value := range m {  
    // order of key sequence different each time  
}
```

- If you require a stable iteration order you must maintain a separate data structure that specifies that order.
 - This example uses a separate sorted slice of keys to print a *map[int]string* in key order:

```
package main  
import (  
    "fmt"  
    "sort"  
)  
  
func main() {  
    var m map[int]string  
    var keys []int  
    for k := range m {  
        keys = append(keys, k)  
    }  
    sort.Ints(keys)  
    for _, k := range keys {  
        fmt.Println("Key:", k, "Value:", m[k])  
    }  
}
```

Map literals

- ✓ Map literals are *like struct literals*, but the keys are required:

```
type Vertex struct {  
    Lat, Long float64  
}  
  
var m = map[string]Vertex {  
    "Bell Labs": Vertex {  
        40.68433, -74.39967,  
    },  
    "Google": Vertex {  
        37.42202, -122.08408,  
    },  
}
```

- ✓ If the top-level type is just a type name, you can omit it from the elements of the literal:

```
type Vertex struct {  
    Lat, Long float64  
}  
  
var m = map[string]Vertex {  
    "Bell Labs": {40.68433, -74.39967},  
    "Google":    {37.42202, -122.08408},  
}
```

- ✓ Same syntax can be used to initialise an *empty map*.
 - This functionally identical to using the `make()` function: `m := map[string]int{}`

Mutating maps

- ✓ Insert or update an element in map `m`: `m[key] = elem`
- ✓ Retrieve an element.
 - If the requested key doesn't exist, we get the value type's zero value: `elem := m[key]`
- ✓ Delete an element.
 - The `delete()` function doesn't return anything, and will do nothing if:
 - the specified key doesn't exist;
 - the map is `nil`.
 - Example: `delete(m, key)`
- ✓ Test that a key is present with a two-value assignment:
 - If key is in `m`, `ok` is true. If not, `ok` is false.
 - If key is not in the map, then `elem` is the zero value for the map's element type.
 - Example: `elem, ok := m[key]`
- ✓ Check for number of items inside map: `n := len(m)`

Common maps operations

✓ Common map operations:

```
make(map[K]V)
len(m)
m[k]
delete(m, k)
```

✓ Maps are not safe for concurrent use!

- One common way to protect maps is with `sync.RWMutex`.
 - This statement declares a 'counter' variable that is an anonymous struct containing a map and an embedded `sync.RWMutex`:

```
var counter = struct {
    sync.RWMutex
    m map[string]int
}{m: make(map[string]int)}
```

- To read from the counter, take the read lock:

```
counter.RLock()
n := counter.m["some_key"]
counter.RUnlock()
fmt.Println("some_key:", n)
```

- To write to the counter, take the write lock:

```
counter.Lock()
counter.m["some_key"]++
counter.Unlock()
```

Slices in Go

Slices general info


- ✓ A slice is a descriptor for a contiguous segment of an *underlying array* and provides access to a numbered sequence of elements **from that array**.
 - The elements can be addressed by *integer* indices **0** through **len(s)-1**.
 - A slice, once initialized, is always associated with an underlying array that holds its elements.
 - A slice therefore shares storage with its *array* and with *other slices of the same array*.
 - By contrast, *distinct arrays always represent distinct storage*.
- ✓ Assigning a slice copies the descriptor, **NOT** the underlying array!
- ✓ Slices may **grow** and **shrink**.
- ✓ The zero value of a slice is nil.
 - A nil slice has a length and capacity of 0.

Slices internals

- ✓ Slices has *length* and *capacity*.
 - Internally, slices looks like this:



- For slice created by `make([]byte, 5)` is structured like this:



Slices under the hood

▶ Check the `runtime.slice` struct definition.

```
15 type slice struct {  
16     > array unsafe.Pointer  
17     > len  int  
18     > cap int  
19 }
```

Slices creation

- ✓ Slices are created with the `make()` function.
 - It works by allocating a zeroed array and returning a slice that refers to that array.
 - A slice created with `make()` always allocates a new, hidden array to which the returned slice value refers.
 - Example:

```
a := make([]int, 5) // Len(a)=5, cap(a)=5
//To specify a capacity, pass a third argument to make:
b := make([]int, 0, 5) // Len(b)=0, cap(b)=5

b = b[:cap(b)] // Len(b)=5, cap(b)=5
b = b[1:]      // Len(b)=4, cap(b)=4
```

- This two forms are *equivalent*:

```
make([]int, 50, 100)
new([100]int)[0:50]
```

Multidimensional Slices (slices of slices)

- ✓ There is no multidimensional slices, instead **slice of slices** can be created:
 - Example with `make()` function:

```
func main() {  
    ss := make([][]uint8, 2) // ss variable is [][]uint8  
    fmt.Printf("ss: %T %v %d/%d\n", ss, ss, len(ss), cap(ss))  
    for i, s := range ss { // s is []uint8  
        s = make([]uint8, 10, 20) // initialize internal slices!  
        fmt.Printf("ss[%d]: %T %v %d/%d\n", i, s, s, len(s), cap(s))  
    }  
}  
  
// OUTPUT:  
//ss: [][]uint8 [][] 2/2  
//ss[0]: []uint8 [0 0 0 0 0 0 0 0 0 0] 10/20  
//ss[1]: []uint8 [0 0 0 0 0 0 0 0 0 0] 10/20
```

Slices expression (short)

- ✓ Slice expressions construct a substring or slice from:
 - string
 - array
 - pointer to array
 - slice.
- ✓ For a string, array, pointer to array, or slice this expression constructs a substring or slice: `a[low : high]`
 - The result has *indices* starting at 0 and *length* equal to `high - low`.
- ✓ For convenience, any of the indices may be omitted.
 - A missing low index defaults to **zero**;
 - A missing high index defaults to the **length of the sliced operand**:

```
a[2:] // same as a[2 : len(a)]  
a[:3] // same as a[0 : 3]  
a[:]  // same as a[0 : len(a)]
```

- ✓ A *constant index* must be non-negative and representable by a value of type int.
 - If the indices are out of range at run time, a **run-time panic** occurs.

Slices expression (full)

✓ Full slice expressions (*rarely used*).

- For an *array*, *pointer to array*, or *slice* *a* (but not a string), the primary expression `a[low : high : max]` constructs a slice of the *same type*, and with the *same length* and *elements* as the simple slice expression `a[low : high]`.
- Additionally, it controls the resulting slice's **capacity** by setting it to `max - low`.
 - *Only the first index may be omitted* (it defaults to 0).
- Useful idea behind using full slice expression is to trigger a new buffer allocation instead of reusing data in original slice.
- Example:

```
a := [5]int{1, 2, 3, 4, 5}
t := a[1:3:5]
// The slice t has type []int, length 2, capacity 4, and elements:
// t[0] == 2
// t[1] == 3
```

Iterating over Slices

- ✓ Slices can be iterated in **for** loop.
- ✓ When *ranging over a slice*, two values are returned for each iteration:
 - the first is the index
 - the second is a copy of the element at that index.
- We can *omit unnecessary variable* during iteration using underscore **_** instead of variable name.

```
package main
import "fmt"

var pow = []int{1, 2, 4, 8, 16, 32, 64, 128}

func main() {
    for i, v := range pow {
        fmt.Printf("2**%d = %d\n", i, v)
    }
}

// OUTPUT:
// 2**0 = 1
// 2**1 = 2
// 2**2 = 4
// 2**3 = 8
// 2**4 = 16
// 2**5 = 32
// 2**6 = 64
// 2**7 = 128
```

Built-in append() function

✓ Function `append()`:

- The variadic function `append()` appends zero or more values `x` to `s` of type `S`, which must be a slice type, and returns the resulting slice, also of type `S`.
- Remember that `append()` may or may not allocate a new slice!
- As a special case, `append()` also accepts a first argument assignable to type `[]byte` with a second argument of `string` type followed by ...
 - This form appends the bytes of the string.

`append(s S, x ...T) S` // *T is the element type of S*

- If the capacity of `s` is not large enough to fit the additional values, `append` allocates a new, sufficiently large underlying array that fits both the existing slice elements and the additional values.
 - Otherwise, `append` re-uses the underlying array.

```
s0 := []int{0, 0}
s1 := append(s0, 2)           // append a single element
s2 := append(s1, 3, 5, 7)     // append multiple elements
s3 := append(s2, s0...)       // append a slice
s4 := append(s3[3:6], s3[2:]...) // append overlapping slice

s1 == []int{0, 0, 2}
s2 == []int{0, 0, 2, 3, 5, 7}
s3 == []int{0, 0, 2, 3, 5, 7, 0, 0}
s4 == []int{3, 5, 7, 2, 3, 5, 7, 0, 0}
```



```
var t []interface{}
t = append(t, 42, 3.1415, "foo") //
t == []interface{}{42, 3.1415, "foo"}
```



```
var b []byte
b = append(b, "bar"...)          // append string contents
b == []byte{'b', 'a', 'r' }
```

Built-in copy() function

✓ Function `copy()`:

- The function `copy` copies slice elements from a source `src` to a destination `dst` and returns the number of elements copied.
- Both arguments must have identical element type `T` and must be assignable to a slice of type `[]T`.
- The number of elements copied is the minimum of `len(src)` and `len(dst)`.
- As a special case, `copy` also accepts a destination argument assignable to type `[]byte` with a source argument of a `string` type.
 - This form copies the bytes from the string into the byte slice:

```
copy(dst, src []T) int
copy(dst []byte, src string) int
```

- Examples:

```
var a = [...]int{0, 1, 2, 3, 4, 5, 6, 7}
var s = make([]int, 6)
var b = make([]byte, 5)
n1 := copy(s, a[0:])           // n1 == 6, s == []int{0, 1, 2, 3, 4, 5}
n2 := copy(s, s[2:])           // n2 == 4, s == []int{2, 3, 4, 5, 4, 5}
n3 := copy(b, "Hello, World!") // n3 == 5, b == []byte("Hello")
```

To be reviewed on the next session

▶ Common pitfalls with slices

▶ Slicing slices pitfalls

Homework

▶ Read: "Go maps in action" by Andrew Gerrand (<https://go.dev/blog/maps>)

▶ Read: "If a map isn't a reference variable, what is it?" by Dave Cheney (<https://dave.cheney.net/2017/04/30/if-a-map-isnt-a-reference-variable-what-is-it>)

▶ Read: "Go Slices: usage and internals" by Andrew Gerrand (<https://blog.golang.org/go-slices-usage-and-internals>)

▶ Read: "Go slices are not dynamic arrays" by Applied Go (<https://appliedgo.net/slices/>)

Next time...

Session08: Struct types and Interfaces in Go

- Types vs. Interfaces: when and what to use?
- Struct types
 - overview of struct types
 - struct embedding
 - struct tags
 - pitfalls with struct types
- Interface types
 - Internals of interface types
 - Well-known pitfalls
 - Type assertions

Thank you

Golang course by Exadel

24 Oct 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

skovtunenکو@exadel.com (mailto:skovtunenکو@exadel.com)

<https://github.com/skovtunenکو> (https://github.com/skovtunenکو)

[@realSKovtunenکو](http://twitter.com/realSKovtunenکو) (http://twitter.com/realSKovtunenکو)

