

# Goroutines and Channels in Go

Session 15

Golang course by Exadel

31 Nov 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

# Agenda

- ▶ Goroutines
- ▶ Channels: bi-directional and uni-directional, buffered and unbuffered
- ▶ Channels Axioms
- ▶ Select statement
- ▶ Blocking vs. non-blocking flows
- ▶ Rules for using channels
- ▶ Next time...

## Before everything else...

▶ You can write production-grade applications in Go without knowing much about:

- goroutines
- channels
- mutexes, etc.

▶ We need goroutines to **better utilize** our hardware!

# Goroutines

# Goroutines 💣

- ▶ The go statement launches a function call as a goroutine

```
go f()  
go f(x, y, ...)
```

- ▶ A goroutine runs **concurrently** (but not necessarily in **parallel**)

- ▶ A goroutine is a **thread of control within the program**, with its **own local variables and stack**.

- ▶ We have no goroutine **ID** ! No "handler" associated with goroutine!

- ▶ *Much cheaper to create and schedule* than operating system threads.

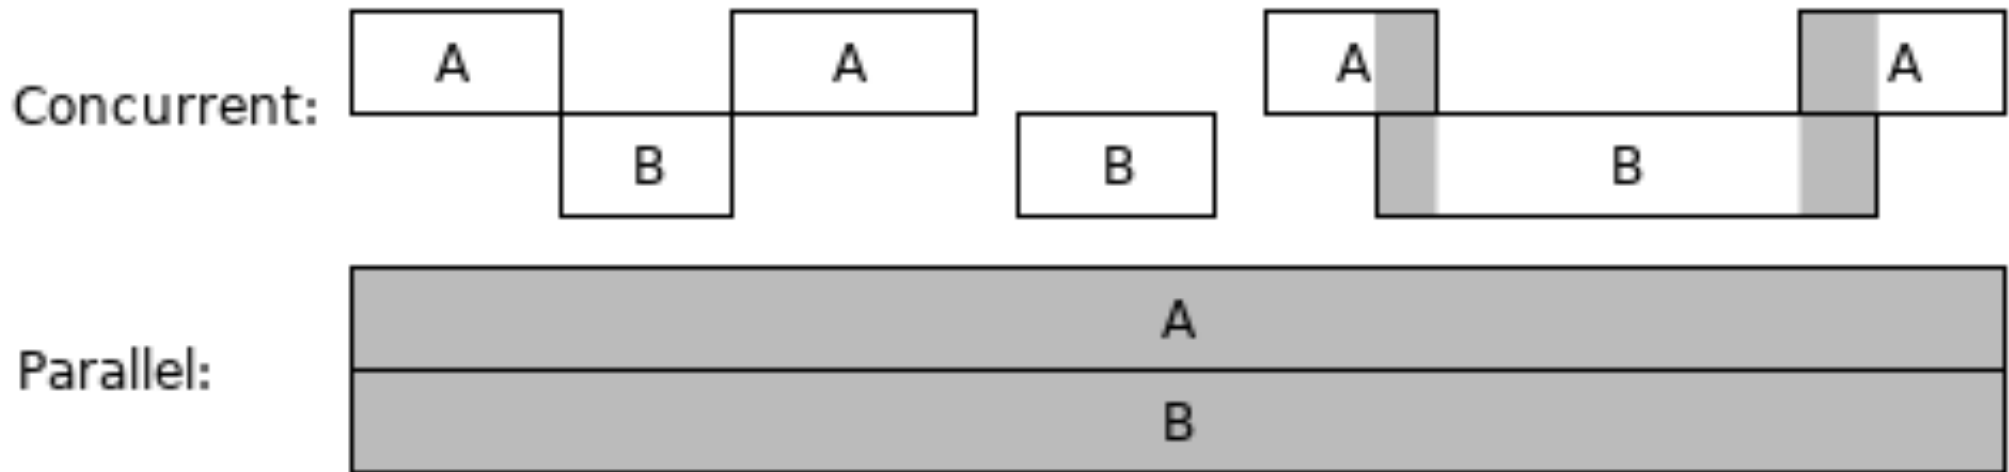
- ▶ **Main** goroutine vs others.

- ▶ Example: `code/01_goroutines_test.go`

Source: "Introduction to Go" by Dave Cheney (<https://github.com/davecheney/introduction-to-go>)

# Understanding Concurrent vs. Parallel Execution

- ▶ Concurrency is about dealing with lots of things at once
- ▶ Parallelism is about doing lots of things at once
- ▶ Concurrency is about structure, parallelism is about execution



Source: <https://talks.golang.org/2012/waza.slide#8> (<https://talks.golang.org/2012/waza.slide#8>)

# Goroutines recap

▶ Goroutines are user-space **lightweight** threads.

- Go runtime **scheduler** will manage them, not the OS. The scheduler is of **work-stealing** type.

▶ Very **lightweight**!

- Owns a **small stack** for local variables: around 2k per 1 goroutine + **stack can grow** as needed
- Faster to create and do context switch
- You can have millions of them in your application

▶ Goroutines are multiplexed onto OS threads

- We can limit the number of OS threads for program written in Go using GOMAXPROCS env var
- `runtime.GOMAXPROCS(32)`

## `go` statement (1/2)

- ✓ A "go" statement starts the execution of a function call as an independent concurrent thread of control, or **goroutine**, within the same address space.
- ✓ The expression must be a function or method call; it cannot be parenthesized.
  - Calls of built-in functions are restricted as for expression statements.
- ✓ The function *value and parameters* are evaluated as usual in the calling goroutine, but unlike with a regular call, program execution does not wait for the invoked function to complete.
  - Instead, the function begins executing independently in a *new goroutine*.
  - When the function terminates, its goroutine also *terminates*.
  - If the function has any return values, they are discarded when the function completes.
- ✓ Example:

```
go Server()  
go func(ch chan<- bool) { for { sleep(10); ch <- true }} (c)
```



## `go` statement (2/2)

✓ You don't need "goroutine ID" at all.

- In Go there is NO goroutine exposed.
- Potential problems include:
  - Because of the high potential for abuse, there is no *<standard>* way to access an identifier for the current goroutine in Go.
  - This may seem draconian, but this actually preserves an important property of the G package ecosystem: *it does not matter if you start a new goroutine to do something.*

```
// That is, for any method of function F:
```

```
F()
```

```
// is almost exactly equivalent to:
```

```
done := make(chan struct{})
```

```
go func() {
```

```
    defer close(done)
```

```
    F()
```

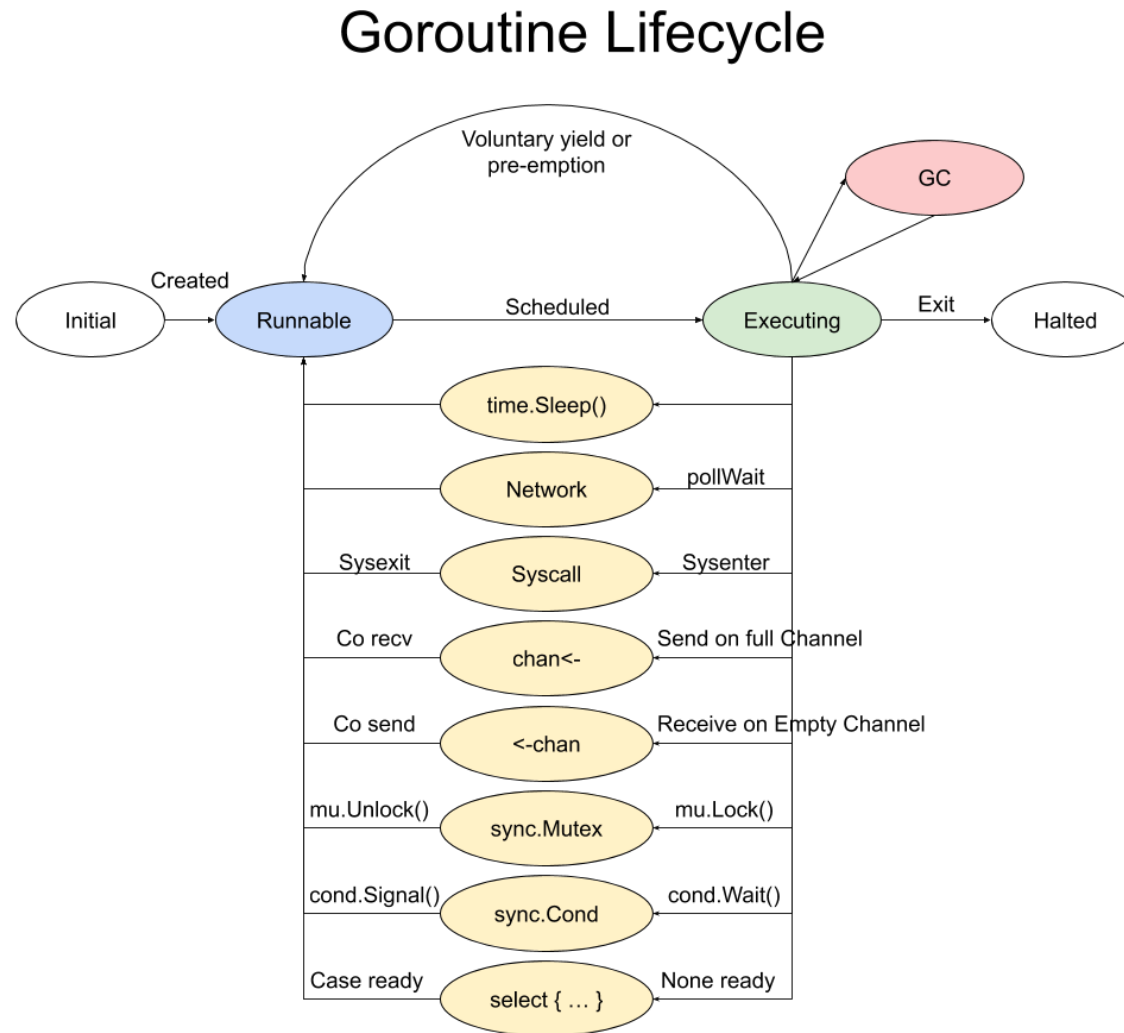
```
}
```

```
<-done
```

✓ When you spawn goroutines, make it clear when - or whether - they exit.

- Try to keep concurrent code simple enough that goroutine lifetimes are obvious.
- If that just isn't feasible, document when and why the goroutines exit.

# Goroutine lifecycle



Source: ["cmd/trace: problems and proposed improvements #33322"](#)



# Goroutine context switching

▶ To understand context switching there is terminology: **M/G/P** :

▶ **G** - represents a user space **Goroutine**

- it has its own stack and blocking info

▶ **P** - represents a logical entity for available **Processors** or a scheduling context.

- The number of P is pre-decided (GOMAXPROCS) and *fixed during the run*.

▶ **M** - stands for **Machines**, they are a representation for OS threads (like POSIX thread)

- **M** will run a set of **G**'s
- Every **Machine** (M) needs a **P** to run **G**.

▶ The same names M/G/P is in use in internal Go's source code.

# Understand Golang scheduler

▶ Awesome 30-mins video: [GopherCon 2018: Kavya Joshi - The Scheduler Saga](https://www.youtube.com/watch?v=YHRO5WQGh0k)

(<https://www.youtube.com/watch?v=YHRO5WQGh0k>)

▶ Another awesome video: [Dmitry Vyukov — Go scheduler: Implementing language with lightweight concurrency](https://www.youtube.com/watch?v=K11rY57K7k) (<https://www.youtube.com/watch?v=K11rY57K7k>)

▶ Watch them! :)

▶ Also useful article: [Go Runtime Scheduler Design Internals](https://freethreads.wordpress.com/2019/01/23/go-runtime-scheduler-design-internals/) (<https://freethreads.wordpress.com/2019/01/23/go-runtime-scheduler-design-internals/>)

# Channels

# Channels: basic info

- ✓ A `channel` provides a mechanism for concurrently executing functions to communicate by sending and receiving values of a specified element type.
  - The value of an uninitialized channel is `nil`.
    - A `nil` channel is never ready for communication.
- ✓ A single channel may be used in:
  - send statements OR
  - receive operations OR
  - calls to the built-in functions `cap()` and `len()` by any number of goroutines without further synchronization.
  - Channels act as first-in-first-out queues.
    - For example, if one goroutine *sends* values on a channel and a second goroutine *receives* them, the values are received *in the order sent*.

# Channels: channel direction

- ✓ The optional `<-` operator specifies the **channel direction**, *send* or *receive*.
  - If *no direction* is given, the channel is **bidirectional**.
  - A channel may be *constrained* only to send or only to receive by conversion or assignment

```
chan T           // can be used to send and receive values of type T
chan<- float64   // can only be used to send float64s
<-chan int       // can only be used to receive ints
```

- ✓ The `<-` operator associates with the leftmost `chan` possible:

```
chan<- chan int   // same as chan<- (chan int)
chan<- <-chan int // same as chan<- (<-chan int)
<-chan <-chan int // same as <-chan (<-chan int)
chan (<-chan int)
```

- ✓ A new, *initialized channel* value can be made using the built-in function `make()`, which takes the channel type and an **optional capacity** as arguments:


```
unbufferedChannel := make(chan int)
bufferedChannel := make(chan int, 100)
```



# Channels: buffered or non-buffered


- ✓ Channels can be **buffered** or **non-buffered**.
- ✓ Unbuffered channels.
  - If *channel capacity* is zero or absent, the channel is unbuffered.
  - The communication succeeds only when both a *sender* and *receiver* are ready.
- ✓ Buffered channels.
  - Channels can be buffered. Provide the buffer length as the second argument to `make()` to initialize a buffered channel.
  - Sends to a buffered channel *block* only when the buffer is full.
  - Receives *block* when the buffer is empty.
- ✓ A **channel** may be *closed* with the built-in function `close()`.
  - The multi-valued assignment form of the receive operator reports whether a received value was sent before the channel was closed.

# Visualization of non-buffered channels



Source: "stackoverflow: What are channels used for?" (<https://stackoverflow.com/a/39831976>)

# Visualization of buffered channels



Source: "stackoverflow: What are channels used for?" (<https://stackoverflow.com/a/39831976>)

# Closing a channel

- ✓ For a channel `c`, the built-in function `close(c)` records that no more values will be sent on the channel.
  - It is an error if `c` is a *receive-only channel*.
- ✓ **Sending to or closing a *closed* channel causes a run-time panic.**
- ✓ Closing the `nil` channel also causes a run-time panic.
- ✓ It is completely normal to not to close channels manually.
  - They will be automatically garbage collected once related (to channels) goroutines finishes.

# Best Practice: Formal arguments for goroutine functions

- ✓ Prefer using formal arguments for the channels you pass to go-routines instead of accessing channels in global scope.
  - You can get *more compiler checking* this way, and *better modularity* too.
  - Example:

```
func main() {  
    c := make(chan string)  
  
    for i := 1; i <= 2; i++ {  
        go func(i int, co chan<- string) {  
            for j := 1; j <= 2; j++ {  
                co <- fmt.Sprintf("hi from %d.%d", i, j)  
            }  
        }(i, c)  
    }  
  
    for i := 1; i <= 4; i++ {  
        fmt.Println(<-c)  
    }  
}  
  
// OUTPUT:  
// Hi from 2.1  
// Hi from 2.2  
// Hi from 1.1  
// Hi from 1.2
```

Source: "Multiple goroutines listening on one channel" [StackOverflow](#)



# Best Practice: Avoid both reading and writing on the same channel

- ✓ **Avoid both reading and writing on the same channel** in a particular go-routine (including the 'main' one).
  - Otherwise, deadlock is a much greater risk.
- ✓ It is generally a good principle to *view buffering as a performance enhancer only*.
  - If your program does not deadlock without buffers, it won't deadlock with buffers either (but the converse is not always true).
  - So, as another rule of thumb, start without buffering then add it later as needed.

# Channel Axioms: send to a nil channel blocks forever

- ✓ A send to a `nil` channel blocks forever:

```
package main

func main() {
    var c chan string
    c <- "let's get started" // deadlock
}
```

Source: "Channel Axioms" by Dave Cheney (<https://dave.cheney.net/2014/03/19/channel-axioms>)



# Channel Axioms: receive from a nil channel blocks forever

- ✓ A receive from a `nil` channel blocks forever:

```
package main
import "fmt"

func main() {
    var c chan string
    fmt.Println(<-c) // deadlock
}
```

- Reason why that happened:
  - If the channel is not initialised then its *buffer size will be zero*.
  - If the size of the channel's buffer is zero, then the channel is *unbuffered*.
  - If the channel is unbuffered, then a send will block until another goroutine is ready to receive.
  - If the channel is `nil` then the sender and receiver have no reference to each other; they are both blocked waiting on independent channels and will never unblock.

Source: "Channel Axioms" by Dave Cheney (<https://dave.cheney.net/2014/03/19/channel-axioms>)

# Channel Axioms: send to a closed channel panics

- ✓ A send to a **closed channel** *panics*:

```
package main
import "fmt"

func main() {
    var c = make(chan int, 100)
    for i := 0; i < 10; i++ {
        go func() {
            for j := 0; j < 10; j++ {
                c <- j
            }
            close(c)
        }()
    }
    for i := range c {
        fmt.Println(i)
    }
}
```

Source: "Channel Axioms" by Dave Cheney (<https://dave.cheney.net/2014/03/19/channel-axioms>)

# Channel Axioms: receive from a closed channel returns the zero value immediately

- ✓ A receive from a **closed channel** returns the *zero value immediately*:

```
package main
import "fmt"

func main() {
    c := make(chan int, 3)
    c <- 1
    c <- 2
    c <- 3
    close(c)
    for i := 0; i < 4; i++ {
        fmt.Printf("%d ", <-c) // prints 1 2 3 0
    }
}
```

Source: "Channel Axioms" by Dave Cheney (<https://dave.cheney.net/2014/03/19/channel-axioms>)

**Next time...**

## We will continue next time

▶ We didn't discuss today:

- select statement
- goroutine gotcha's

## Homework:

▶ "3 Golang Channel Red Flags" (<https://hackmongo.com/post/three-golang-channel-red-flags/>)

▶ Must read: [Channel Axioms](https://dave.cheney.net/2014/03/19/channel-axioms) (<https://dave.cheney.net/2014/03/19/channel-axioms>)

▶ [Channels in Go](https://go101.org/article/channel.html) (<https://go101.org/article/channel.html>)

▶ [Principles of designing Go APIs with channels](https://inconshreveable.com/07-08-2014/principles-of-designing-go-apis-with-channels/) (<https://inconshreveable.com/07-08-2014/principles-of-designing-go-apis-with-channels/>)

▶ [Understanding real-world concurrency bugs in Go](https://blog.acolyer.org/2019/05/17/understanding-real-world-concurrency-bugs-in-go/) (<https://blog.acolyer.org/2019/05/17/understanding-real-world-concurrency-bugs-in-go/>)

▶ Just for information, not to be used: [A pattern for overcoming non-determinism of Golang select statement](https://medium.com/@pedram.esmaeeli/a-pattern-for-overcoming-non-determinism-of-golang-select-statement-139dbe93db98) (<https://medium.com/@pedram.esmaeeli/a-pattern-for-overcoming-non-determinism-of-golang-select-statement-139dbe93db98>)

## Next session...

 Session16:

**Most useful packages: context, sync**

- Context package
- Sync package

# Thank you

Golang course by Exadel

31 Nov 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

[skovtunenکو@exadel.com](mailto:skovtunenکو@exadel.com) (mailto:skovtunenکو@exadel.com)

<https://github.com/skovtunenکو> (https://github.com/skovtunenکو)

[@realSKovtunenکو](http://twitter.com/realSKovtunenکو) (http://twitter.com/realSKovtunenکو)



