

# Channels - Part 2. The most Useful package in Go: sync

Session 16

Golang course by Exadel

05 Dec 2022

Sergio Kovtunenکو


Lead backend developer, Exadel

# Agenda

- ▶ Send (to channel) statement details
- ▶ Receive (from channel) operator
- ▶ for-range loop with
- ▶ Goroutine gotchas
- ▶ Select statement
- ▶ Blocking vs. non-blocking flows
- ▶ The most useful package: sync
- ▶ Advanced concurrency patterns
- ▶ Next time...

# Goroutines and channels - Part 2

# Send (to channel) statement

- ✓ A **send** statement *sends a value on a channel*.
- ✓ The channel expression must be:
  - of channel type **AND**
  - the channel direction must permit send operations **AND**
  - the type of the value to be sent must be assignable to the channel's element type.
- ✓ Both the *channel* and the value expression are evaluated before communication begins.
- ✓ Communication blocks until the send can proceed.
- ✓ Potential issue with **mutable** data (pointers to data) send into channel. 
- ✓ Restrictions:
  - A **send** on an unbuffered channel can proceed if a receiver is ready.
  - A **send** on a buffered channel can proceed if there is room in the buffer.
  - A **send** on a closed channel proceeds by causing a run-time panic.
  - A **send** on a nil channel blocks forever.

```
ch <- 3 // send value 3 to channel ch
```

# Receive (from channel) operator: recap from the last session

- ✓ For an operand *ch* of *channel* type, the value of the receive operation `<-ch` is the value *received from the channel ch*.
- ✓ The **channel direction** must:
  - permit receive operations
  - and the type of the receive operation is the element type of the channel.
- ✓ The expression blocks until a value is available.
- ✓ Receiving from a **nil** channel **blocks forever**.
- ✓ A **receive operation** on a **closed channel** can always proceed immediately
  - yielding the element type's zero value after any previously sent values have been received.

```
v1 := <-ch
v2 = <-ch
f(<-ch)
<-strobe // wait until clock pulse and discard received value
```

## Receive (from channel) operator: comma-ok check

- ✓ A **receive expression** used in an assignment or initialization of the special form
  - yields an additional untyped boolean result reporting whether the communication succeeded.
  - The value of **ok** is **true** if the value received was delivered by a successful send operation to the channel
    - or **false** if it is a zero value generated because the channel is closed and empty.
  - Example:

```
x, ok = <-ch
x, ok := <-ch
var x, ok = <-ch
var x, ok T = <-ch
```

# For-range loop over channel

- ✓ For `channels`, the iteration values produced are the successive values sent on the channel until the channel is closed.
- ✓ If the `channel` is `nil`, the range expression blocks forever.

```
package main
import "fmt"
func FibonacciProducer(ch chan int, count int) {
    n2, n1 := 0, 1
    for count >= 0 {
        ch <- n2
        count--
        n2, n1 = n1, n2+n1
    }
    close(ch)
}
func main() {
    ch := make(chan int)
    go FibonacciProducer(ch, 10)
    idx := 0
    // To break such iteration channel needs to be closed explicitly.
    // Otherwise range would block forever in the same way as for nil channel.
    for num := range ch {
        fmt.Printf("F(%d): %t%d\n", idx, num)
        idx++
    }
}
```

## For-range loop over channel - example

▶ Example: `code/chanexample/01_chan_iteration_test.go`



# Blocked Goroutines and Resource Leaks: problem statement

## ✓ Blocked Goroutines and Resource Leaks:

- Fetching the first result from a number of targets is one of them (goroutine leak in this example):
  - To avoid the leaks you need to make sure all goroutines exit.

```
func First(query string, replicas ...Search) Result {  
    c := make(chan Result) // Unbuffered (synchronous) channel!  
  
    searchReplica := func(i int) { c <- replicas[i](query) }  
    for i := range replicas {  
        go searchReplica(i)  
    }  
  
    return <-c // This means that only the first goroutine returns.  
    // All other goroutines are stuck trying to send their results.  
    // This means if you have more than one replica each call will leak resources.  
}
```

Source: "50 Shades of Go: Traps, Gotchas, and Common Mistakes for New Golang Devs"

([http://devs.cloudimmunity.com/gotchas-and-common-mistakes-in-go-golang/index.html#blocked\\_goroutines](http://devs.cloudimmunity.com/gotchas-and-common-mistakes-in-go-golang/index.html#blocked_goroutines))

# Blocked Goroutines and Resource Leaks: solution 1

## ▶ Solution 1:

✓ **Solution 1:** Fix using a buffered result channel big enough to hold all results:

```
func First(query string, replicas ...Search) Result {  
    // One potential solution is to use a buffered result channel big enough to hold all results:  
    c := make(chan Result, len(replicas))  
  
    searchReplica := func(i int) { c <- replicas[i](query) }  
    for i := range replicas {  
        go searchReplica(i)  
    }  
    return <-c  
}
```

# Blocked Goroutines and Resource Leaks: solution 2

## ▶ Solution 2:

- ✓ **Solution 2:** Fix using a select statement with a default case and a buffered result channel that can hold one value.
  - The default case ensures that the goroutines don't get stuck even when the result channel can't receive messages:

```
func First(query string, replicas ...Search) Result {  
    c := make(chan Result, 1) // Buffered channel with size 1  
    searchReplica := func(i int) {  
        select {  
        case c <- replicas[i](query):  
        default:  
            // Will be executed for the rest of goroutines blocked on sending values!  
        }  
    }  
    for i := range replicas {  
        go searchReplica(i)  
    }  
    return <-c  
}
```

# Blocked Goroutines and Resource Leaks: solution 3

## ▶ Solution 3:

✓ **Solution 3:** Fix using a special cancellation channel to interrupt the workers:

```
func First(query string, replicas ...Search) Result {
    c := make(chan Result)

    done := make(chan struct{}) // Special cancellation channel
    defer close(done)

    searchReplica := func(i int) {
        select {
        case c <- replicas[i](query):
        case <- done:
            // Will be executed for the rest of goroutines blocked on sending values!
        }
    }

    for i := range replicas {
        go searchReplica(i)
    }

    return <-c
}
```

## `select` statement: general information

- ✓ A "select" statement chooses which of a set of possible send or receive operations will proceed.
  - It looks similar to a "switch" statement but with the cases all referring to communication operations.
- ✓ Since communication on nil channels can never proceed:
  - select with only nil channels and no default case blocks forever.
  - but using nil chans will disable a select case.

# `select` statement: execution flow details

- ✓ Execution of a "select" statement proceeds in several steps:
  - For all the cases in the statement, the channel operands of **receive operations** and the **channel** and **right-hand-side expressions of send statements** are evaluated exactly once, in source order, upon entering the "select" statement.
    - The result is a set of channels to **receive from** or **send to**, and the corresponding values to send.
    - Any side effects in that evaluation will occur irrespective of which (if any) communication operation is selected to proceed.
      - Expressions on the left-hand side of a "Receive Statement" with a short *variable declaration* or *assignment* are not yet evaluated.
  - If one or more of the communications can proceed, a single one that can proceed is chosen via a uniform pseudo-random selection.
    - Otherwise, if there is a default case, that case is chosen.
      - If there is no **default case**, the "select" statement blocks until at least one of the communications can proceed.
  - Unless the selected case is the **default case**, the respective communication operation is executed.
  - If the selected case is a "Receive Statement" (with a short variable declaration or an assignment), the left-hand side expressions are evaluated and the received value (or values) are assigned.
  - The statement list of the selected case is executed.

# `select` statement: example

✓ Example:

```
func demo() {  
  var a []int  
  var c, c1, c2, c3, c4 chan int  
  var i1, i2 int  
  select {  
  case i1 = <-c1:  
    print("received ", i1, " from c1\n")  
  case c2 <- i2:  
    print("sent ", i2, " to c2\n")  
  case i3, ok := (<-c3): // same as: i3, ok := <-c3  
    if ok {  
      print("received ", i3, " from c3\n")  
    } else {  
      print("c3 is closed\n")  
    }  
  case a[f()] = <-c4:  
    // same as:  
    // case t := <-c4  
    //   a[f()] = t  
  default:  
    print("no communication\n")  
  }  
  
  for { // send random sequence of bits to c  
    select {  
    case c <- 0: // note: no statement, no fallthrough, no folding of cases  
    case c <- 1:  
    }  
  }  
  
  select {} // block forever  
}
```

``sync` package`



# Overview of `sync` package

▶ `sync` [package documentation](https://pkg.go.dev/sync) (https://pkg.go.dev/sync)

▶ Package **sync** provides basic synchronization primitives such as mutual exclusion locks.

▶ Other than the **Once** and **WaitGroup** types, *most are intended for use by low-level library routines.*

- *Higher-level synchronization is better done via channels and communication.*
- Values containing the types defined in this package **MUST not be copied.**

## `sync` package: sync.Mutex

▶ Type: [sync.Mutex](https://pkg.go.dev/sync#Mutex) (<https://pkg.go.dev/sync#Mutex>)

▶ It allows a **mutual exclusion** on a shared resource (no simultaneous access):

```
mutex := &sync.Mutex{}
```

```
mutex.Lock()
```

```
mutex.Unlock()
```

▶ A `sync.Mutex` cannot be copied!!!

## `sync` package: sync.RWMutex

▶ Type: [sync.RWMutex](https://pkg.go.dev/sync#RWMutex) (<https://pkg.go.dev/sync#RWMutex>).

▶ API usage example:

```
mutex := &sync.RWMutex{}
```

```
mutex.Lock()
```

```
mutex.Unlock()
```

```
mutex.RLock()
```

```
mutex.RUnlock()
```

▶ A `sync.RWMutex` allows either **at least** one reader or **exactly** one writer whereas a `sync.Mutex` allows **exactly one** reader or writer.

▶ A `sync.RWMutex` should rather be used when we have **frequent reads** and **infrequent** writes.

▶ A `sync.RWMutex` **cannot be copied!!!**

## `sync` package: `sync.WaitGroup`

▶ Type `sync.WaitGroup` (<https://pkg.go.dev/sync#WaitGroup>)

▶ Provides an idiomatic way for a goroutine to wait for the **completion of a collection of goroutines**.

▶ `sync.WaitGroup` holds an internal counter.

- If this counter is equal to 0, the `Wait()` method returns immediately.
- Otherwise, it is blocked until the counter is 0.

```
wg := &sync.WaitGroup{}
for i := 0; i < 8; i++ {
    wg.Add(1)
    go func() {
        // Do something...
        wg.Done()
    }()
}
wg.Wait() // Continue execution afterwards...
```

## `sync` package: sync.Once

▶ To execute a code only once: [sync.Once](https://pkg.go.dev/sync#Once) (https://pkg.go.dev/sync#Once)

▶ Example:

```
once := &sync.Once{}
```

```
for i := 0; i < 4; i++ {  
    i := i
```

```
    go func() {  
        once.Do(func() {  
            fmt.Printf("first %d\n", i)  
        })  
    }()  
}
```

## `sync` package: rest of APIs

▶ Concurrent pool, in charge to hold safely a set of objects: [sync.Pool](https://pkg.go.dev/sync#Pool) (https://pkg.go.dev/sync#Pool).

- `Get() interface{}` to retrieve an element;
- `Put(interface{})` to add an element;

▶ Concurrent map implementation: [sync.Map](https://pkg.go.dev/sync#Map) (https://pkg.go.dev/sync#Map)

▶ The less frequently used primitive: [sync.Cond](https://pkg.go.dev/sync#Cond) (https://pkg.go.dev/sync#Cond)

- It is used to emit a signal (one-to-one) or broadcast a signal (one-to-many) to goroutine(s).

# Advanced concurrency patterns

## General info about advanced patterns: fan-out, fan-in

### ▶ Fan-out:

- **Multiple functions can read from the same channel** until that channel is closed;
- This provides a way to distribute work amongst a group of workers to parallelize CPU use and I/O.

### ▶ Fan-in:

- A function can read from multiple inputs and proceed until all are closed by multiplexing the input channels onto a single channel that's closed when all the inputs are closed.



# Generators

▶ Example code: `code/02_advancedexample/01_generate_data_test.go`

▶ Remember this code:

```
func asChan(nums ...int) <-chan int {  
    ch := make(chan int)  
    go func() {  
        defer close(ch)  
        for _, val := range nums {  
            ch <- val  
        }  
    }()  
    return ch  
}
```

## Merging channels together

▶ Merge two (or N known channels):

- Example code: `code/02_advancedexample/02_merge_two_channels_test.go`

▶ Merge many (unknown) number of channels of same type:

- Example code: `code/02_advancedexample/03_merge_many_channels_test.go`

## Homework:

▶ Recap: [Go Concurrency Patterns: Pipelines and cancellation](https://go.dev/blog/pipelines) (<https://go.dev/blog/pipelines>)

## Next session...

 Session17:

### HTTP servers and routers

- standard package context
- Typical HTTP Server
- HTTP server settings
- Testing HTTP server

# Thank you

Golang course by Exadel

05 Dec 2022

Sergio Kovtunenکو

Lead backend developer, Exadel

[skovtunenکو@exadel.com](mailto:skovtunenکو@exadel.com) (mailto:skovtunenکو@exadel.com)

<https://github.com/skovtunenکو> (https://github.com/skovtunenکو)

[@realSKovtunenکو](http://twitter.com/realSKovtunenکو) (http://twitter.com/realSKovtunenکو)

